

Assignment 5.2: News Classifier

```
In [19]: import numpy as np
import matplotlib.pyplot as plt

from keras import models
from keras import layers
from tensorflow.keras import optimizers
from keras.datasets import reuters
from keras import losses
from keras import metrics
from keras.utils.np_utils import to_categorical
```

```
In [20]: (train_data, train_labels), (test_data, test_labels) = reuters.load_data(num
```

Vectorize Data

```
In [30]: def vectorize_sequences(sequences, dimension=10000):
        results = np.zeros((len(sequences), dimension))
        for i, sequence in enumerate(sequences):
            results[i, sequence] = 1.
        return results
```

```
In [31]: x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
```

```
In [32]: one_hot_train_labels = to_categorical(train_labels)
one_hot_test_labels = to_categorical(test_labels)
```

Building the Network

```
In [33]: model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(46, activation='softmax'))
```

```
In [34]: model.compile(optimizer='rmsprop',
                    loss='categorical_crossentropy',
                    metrics=['accuracy'])
```

Validation Set

```
In [35]: x_val = x_train[:1000]
partial_x_train = x_train[1000:]

y_val = one_hot_train_labels[:1000]
partial_y_train = one_hot_train_labels[1000:]
```

Training the Model

Epoch 1/20
16/16 [=====] - 2s 55ms/step - loss: 2.6948 - accuracy: 0.5563 - val_loss: 1.7709 - val_accuracy: 0.6510
Epoch 2/20
16/16 [=====] - 1s 40ms/step - loss: 1.4326 - accuracy: 0.7117 - val_loss: 1.3039 - val_accuracy: 0.7110
Epoch 3/20
16/16 [=====] - 1s 36ms/step - loss: 1.0663 - accuracy: 0.7745 - val_loss: 1.1308 - val_accuracy: 0.7670
Epoch 4/20
16/16 [=====] - 1s 47ms/step - loss: 0.8533 - accuracy: 0.8206 - val_loss: 1.0289 - val_accuracy: 0.7870
Epoch 5/20
16/16 [=====] - 1s 42ms/step - loss: 0.6873 - accuracy: 0.8563 - val_loss: 0.9745 - val_accuracy: 0.7990
Epoch 6/20
16/16 [=====] - 1s 45ms/step - loss: 0.5519 - accuracy: 0.8854 - val_loss: 0.9256 - val_accuracy: 0.8040
Epoch 7/20
16/16 [=====] - 1s 47ms/step - loss: 0.4496 - accuracy: 0.9047 - val_loss: 0.9409 - val_accuracy: 0.7930
Epoch 8/20
16/16 [=====] - 1s 44ms/step - loss: 0.3626 - accuracy: 0.9230 - val_loss: 0.8769 - val_accuracy: 0.8240
Epoch 9/20
16/16 [=====] - 1s 81ms/step - loss: 0.2989 - accuracy: 0.9360 - val_loss: 0.8974 - val_accuracy: 0.8220
Epoch 10/20
16/16 [=====] - 1s 62ms/step - loss: 0.2510 - accuracy: 0.9426 - val_loss: 0.9231 - val_accuracy: 0.8140
Epoch 11/20
16/16 [=====] - 1s 38ms/step - loss: 0.2150 - accuracy: 0.9496 - val_loss: 0.9047 - val_accuracy: 0.8190
Epoch 12/20
16/16 [=====] - 1s 44ms/step - loss: 0.1908 - accuracy: 0.9501 - val_loss: 0.9395 - val_accuracy: 0.8180
Epoch 13/20
16/16 [=====] - 1s 47ms/step - loss: 0.1714 - accuracy: 0.9528 - val_loss: 0.9483 - val_accuracy: 0.8130
Epoch 14/20
16/16 [=====] - 1s 54ms/step - loss: 0.1537 - accuracy: 0.9531 - val_loss: 0.9705 - val_accuracy: 0.8260
Epoch 15/20
16/16 [=====] - 1s 58ms/step - loss: 0.1444 - accuracy: 0.9557 - val_loss: 0.9949 - val_accuracy: 0.8100
Epoch 16/20
16/16 [=====] - 1s 47ms/step - loss: 0.1347 - accuracy: 0.9560 - val_loss: 1.0125 - val_accuracy: 0.8130
Epoch 17/20
16/16 [=====] - 1s 39ms/step - loss: 0.1274 - accuracy: 0.9564 - val_loss: 1.0403 - val_accuracy: 0.8040
Epoch 18/20
16/16 [=====] - 1s 50ms/step - loss: 0.1237 - accuracy: 0.9546 - val_loss: 1.0747 - val_accuracy: 0.8010
Epoch 19/20
16/16 [=====] - 1s 51ms/step - loss: 0.1187 - accuracy:

```
racy: 0.9567 - val_loss: 1.0842 - val_accuracy: 0.7970
Epoch 20/20
16/16 [=====] - 1s 54ms/step - loss: 0.1136 - accu
racy: 0.9594 - val_loss: 1.1199 - val_accuracy: 0.8000
```

Training and Validation Loss

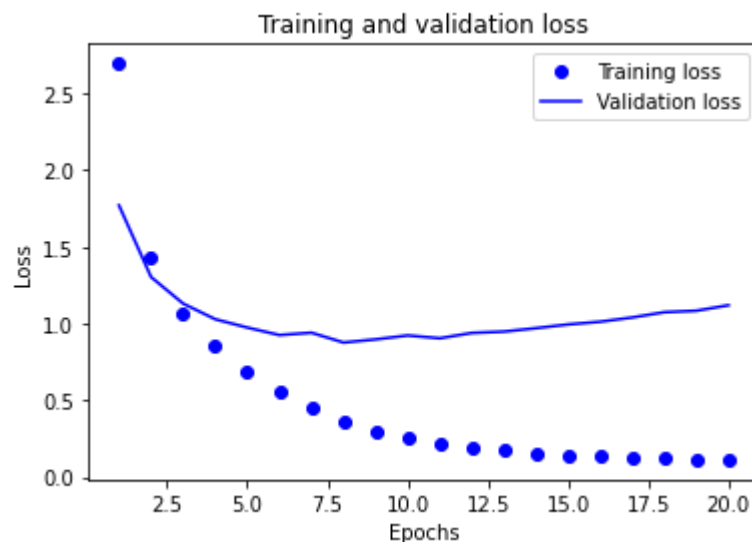
```
In [37]: loss = history.history['loss']
val_loss = history.history['val_loss']
```

```
In [38]: epochs = range(1, 21)
```

```
In [39]: plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')

plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



Training and Validation Accuracy

```
In [40]: plt.clf()
```

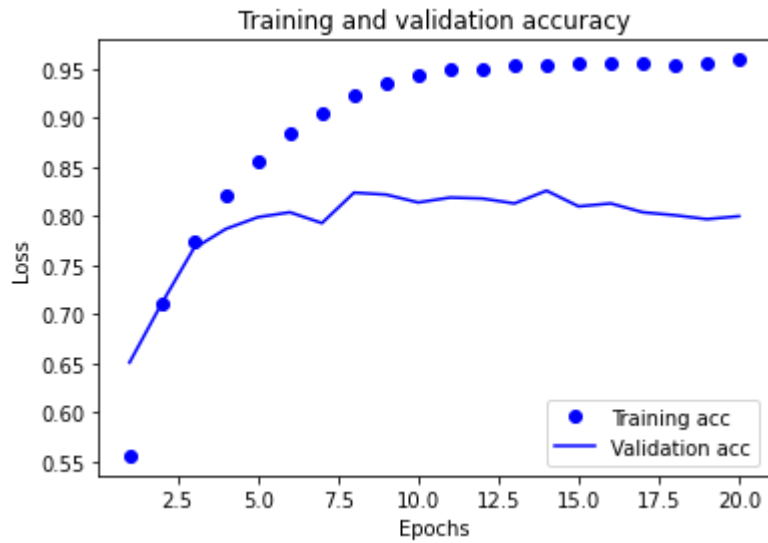
<Figure size 432x288 with 0 Axes>

```
In [43]: acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
```

```
In [44]: plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')

plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
```

```
plt.show()
```



Results

```
In [46]: results = model.evaluate(x_test, one_hot_test_labels)
results
```

71/71 [=====] - 1s 6ms/step - loss: 1.2256 - accuracy: 0.7792

```
Out[46]: [1.2256207466125488, 0.7791629433631897]
```

The model has a 77.92% accuracy.