

Assignment 6.1: MNIST

```
In [27]: import pandas as pd
import matplotlib.pyplot as plt

from keras import layers
from keras import models
from tensorflow.keras.utils import to_categorical

from keras.datasets import mnist
```

Initiating a Convnet

```
In [2]: model = models.Sequential()
model.add(layers.Conv2D(32,(3,3), activation='relu', input_shape=(28, 28, 1))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(64,(3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(64,(3,3), activation='relu'))
```

Adding a Classifier

```
In [4]: model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))
```

```
In [5]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928
flatten (Flatten)	(None, 576)	0
dense (Dense)	(None, 64)	36928
dense_1 (Dense)	(None, 10)	650

=====
Total params: 93,322
Trainable params: 93,322
Non-trainable params: 0
=====

Training Convnet on MNIST Data

```
In [8]: (train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

```
In [9]: train_images = train_images.reshape((60000, 28, 28, 1))
train_images = train_images.astype('float32') / 255

test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype('float32') / 255

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
```

```
In [10]: model.compile(optimizer='rmsprop',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
```

```
In [11]: model.fit(train_images, train_labels, epochs=5, batch_size=64)
```

```

Epoch 1/5
938/938 [=====] - 40s 42ms/step - loss: 0.1747 - a
ccuracy: 0.9446
Epoch 2/5
938/938 [=====] - 46s 49ms/step - loss: 0.0480 - a
ccuracy: 0.9851
Epoch 3/5
938/938 [=====] - 38s 41ms/step - loss: 0.0324 - a
ccuracy: 0.9904
Epoch 4/5
938/938 [=====] - 40s 43ms/step - loss: 0.0246 - a
ccuracy: 0.9921
Epoch 5/5
938/938 [=====] - 43s 46ms/step - loss: 0.0185 - a
ccuracy: 0.9944

```

Out[11]: <keras.callbacks.History at 0x7fec10857cd0>

Saving the Model

```

In [12]: model.save('results/model-6-1')

INFO:tensorflow:Assets written to: results/model-6-1/assets

```

Saving the Predictions

```

In [19]: pred = model.predict(test_images)

pred = pd.DataFrame(pred)
pred.to_csv('results/model-6-1/predictions.csv')

```

Saving the Metrics

```

In [23]: test_loss, test_acc = model.evaluate(test_images, test_labels)

313/313 [=====] - 2s 6ms/step - loss: 0.0331 - acc
uracy: 0.9903

In [24]: print("Test Accuracy: ", test_acc)
print("Test Loss: ", test_loss)

Test Accuracy: 0.9902999997138977
Test Loss: 0.03312191739678383

In [25]: lines = ['Test Accuracy: 0.9902999997138977', 'Test Loss: 0.03312191739678383']

In [26]: with open('results/model-6-1/metrics.txt', 'w') as f:
f.write('\n'.join(lines))

```

Saving the Validation Plots

```

In [28]: history = model.fit(train_images, train_labels, epochs=5, batch_size=64)

```

```
Epoch 1/5
938/938 [=====] - 46s 49ms/step - loss: 0.0151 - a
ccuracy: 0.9953
Epoch 2/5
938/938 [=====] - 42s 44ms/step - loss: 0.0124 - a
ccuracy: 0.9963
Epoch 3/5
938/938 [=====] - 40s 43ms/step - loss: 0.0102 - a
ccuracy: 0.9970
Epoch 4/5
938/938 [=====] - 37s 39ms/step - loss: 0.0093 - a
ccuracy: 0.9972
Epoch 5/5
938/938 [=====] - 38s 40ms/step - loss: 0.0073 - a
ccuracy: 0.9977
```

```
In [34]: acc = history.history['accuracy']
loss = history.history['loss']
```

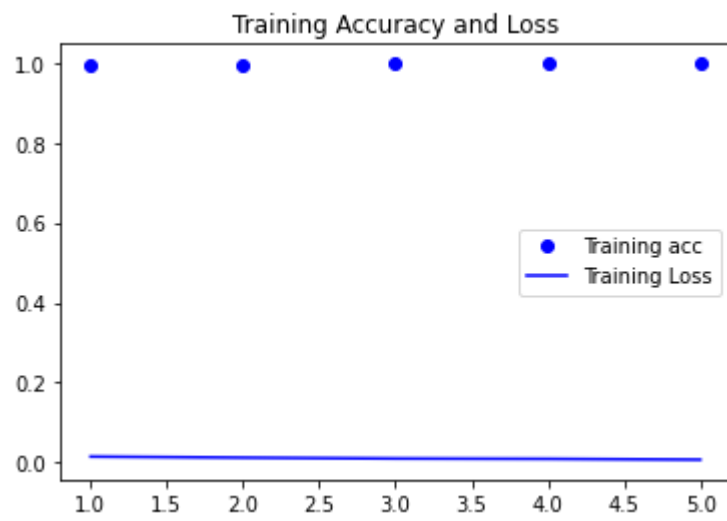
```
In [40]: history.history
```

```
Out[40]: {'loss': [0.01508525013923645,
0.012374487705528736,
0.010205096565186977,
0.009319016709923744,
0.007304158993065357],
'accuracy': [0.9953166842460632,
0.9962666630744934,
0.9969666600227356,
0.9972166419029236,
0.9977499842643738]}
```

```
In [35]: epochs = range(1, len(acc) + 1)
```

```
In [37]: plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, loss, 'b', label='Training Loss')
plt.title('Training Accuracy and Loss')
plt.legend()

plt.show()
```



```
In [41]: plt.savefig('results/model-6-1/validationplot.jpg')
```

<Figure size 432x288 with 0 Axes>