

## Assignment 6.2.a: CIFAR10

```
In [ ]: import sys
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt

from keras.datasets import cifar10
from keras.models import Sequential

from keras.layers import Conv2D
from keras.layers import MaxPooling2D
from keras.layers import Dense
from keras.layers import Flatten
from keras.layers import Dropout

from tensorflow.keras.utils import to_categorical
from tensorflow.keras.optimizers import SGD

from keras.preprocessing import image
from keras.preprocessing.image import ImageDataGenerator

from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
```

### Import Data

```
In [3]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

```
In [4]: y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
```

### Format Photos

```
In [5]: train_norm = x_train.astype('float32')
test_norm = x_test.astype('float32')

train_norm = train_norm / 255.0
test_norm = test_norm / 255.0
```

### Initiating a Model

```
In [6]: model = Sequential()
model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))
model.add(Dense(10, activation='softmax'))
```

```
In [7]: opt = SGD(learning_rate=0.001, momentum=0.9)
```

```
In [8]: model.compile(optimizer=opt,  
                    loss='categorical_crossentropy',  
                    metrics=['accuracy'])
```

```
In [9]: history = model.fit(x_train, y_train, epochs=5, batch_size=64, validation_data=(x_test, y_test))  
  
Epoch 1/5  
782/782 [=====] - 120s 152ms/step - loss: 17648.5059 - accuracy: 0.0986 - val_loss: 2.3026 - val_accuracy: 0.1000  
Epoch 2/5  
782/782 [=====] - 121s 155ms/step - loss: 2.3027 - accuracy: 0.0985 - val_loss: 2.3026 - val_accuracy: 0.1000  
Epoch 3/5  
782/782 [=====] - 140s 179ms/step - loss: 2.3026 - accuracy: 0.0974 - val_loss: 2.3026 - val_accuracy: 0.1000  
Epoch 4/5  
782/782 [=====] - 131s 167ms/step - loss: 2.3026 - accuracy: 0.0986 - val_loss: 2.3026 - val_accuracy: 0.1000  
Epoch 5/5  
782/782 [=====] - 116s 148ms/step - loss: 2.3026 - accuracy: 0.0965 - val_loss: 2.3026 - val_accuracy: 0.1000
```

### Saving the Model

```
In [10]: model.save('results/model-6-2-a')  
  
INFO:tensorflow:Assets written to: results/model-6-2-a/assets
```

### Saving the Predictions

```
In [11]: pred = model.predict(x_test)  
  
pred = pd.DataFrame(pred)  
pred.to_csv('results/model-6-2-a/predictions.csv')
```

### Saving the Metrics

```
In [12]: test_loss, test_acc = model.evaluate(x_test, y_test)  
  
313/313 [=====] - 5s 15ms/step - loss: 2.3026 - accuracy: 0.1000
```

```
In [13]: print("Test Accuracy: ", test_acc)  
        print("Test Loss: ", test_loss)  
  
Test Accuracy:  0.10000000149011612  
Test Loss:    2.302586078643799
```

```
In [14]: lines = ['Test Accuracy:  0.10000000149011612', 'Test Loss:    2.302586078643799']
```


```
In [15]: with open('results/model-6-2-a/metrics.txt', 'w') as f:  
        f.write('\n'.join(lines))
```


## Saving the Validation Plots

```
In [16]: acc = history.history['accuracy']  
val_acc = history.history['val_accuracy']  
  
loss = history.history['loss']  
val_loss = history.history['val_loss']
```

```
In [17]: epochs = range(1, len(acc) + 1)
```

```
In [18]: plt.plot(epochs, acc, 'bo', label='Training acc')  
plt.plot(epochs, val_acc, 'b', label='Validation acc')  
plt.title('Training and validation accuracy')  
plt.legend()  
  
plt.figure()  
  
plt.plot(epochs, loss, 'bo', label='Training loss')  
plt.plot(epochs, val_loss, 'b', label='Validation loss')  
plt.title('Training and validation loss')  
  
plt.legend()  
  
plt.show()
```





```
In [19]: plt.savefig('results/model-6-2-a/validationplot.jpg')
```

<Figure size 432x288 with 0 Axes>

## Assignment 6.2.b: CIFAR10 with Dropout & Augmentation

### Initiating the Model

```
In [20]: modelB = Sequential()
modelB.add(Conv2D(32, kernel_size=(3,3), activation='relu', input_shape=(32, 32, 3)))
modelB.add(MaxPooling2D(pool_size=(2, 2)))
modelB.add(Flatten())
modelB.add(Dense(16, activation='relu'))
modelB.add(Dense(1, activation='sigmoid'))
```

```
In [21]: modelB.compile(loss='binary_crossentropy',
                        optimizer='adam',
                        metrics=['accuracy'])
```

### Image Augmentation

```
In [22]: train_datagen = ImageDataGenerator(
        rotation_range=40,
        width_shift_range=0.2,
        height_shift_range=0.2,
        shear_range=0.2,
        zoom_range=0.2,
        horizontal_flip=True,
        fill_mode='nearest')

test_datagen = ImageDataGenerator(rescale=1./255)
```

```
In [23]: train_generator = train_datagen.flow(x_train, y_train, batch_size=32)
```

```
In [24]: validation_generator = train_datagen.flow(x_train, y_train, batch_size=32)
```

```
In [25]: test_generator = test_datagen.flow(x_test, batch_size=1)
```

### Fitting the Model

```
In [26]: historyB = modelB.fit(
        train_generator,
        steps_per_epoch=100,
        epochs=10,
        validation_data=validation_generator,
        validation_steps=50)
```

```
Epoch 1/10
100/100 [=====] - 9s 85ms/step - loss: 1.9452 - ac
curacy: 0.8600 - val_loss: 0.3813 - val_accuracy: 0.8960
Epoch 2/10
100/100 [=====] - 7s 71ms/step - loss: 0.3712 - ac
curacy: 0.8968 - val_loss: 0.3678 - val_accuracy: 0.8960
Epoch 3/10
100/100 [=====] - 10s 102ms/step - loss: 0.3631 -
accuracy: 0.8983 - val_loss: 0.3594 - val_accuracy: 0.9000
Epoch 4/10
100/100 [=====] - 11s 106ms/step - loss: 0.3545 -
accuracy: 0.8985 - val_loss: 0.3516 - val_accuracy: 0.9000
Epoch 5/10
100/100 [=====] - 14s 139ms/step - loss: 0.3468 -
accuracy: 0.9000 - val_loss: 0.3454 - val_accuracy: 0.9000
Epoch 6/10
100/100 [=====] - 8s 80ms/step - loss: 0.3427 - ac
curacy: 0.9000 - val_loss: 0.3382 - val_accuracy: 0.9000
Epoch 7/10
100/100 [=====] - 7s 72ms/step - loss: 0.3402 - ac
curacy: 0.8998 - val_loss: 0.3399 - val_accuracy: 0.8995
Epoch 8/10
100/100 [=====] - 5s 54ms/step - loss: 0.3387 - ac
curacy: 0.8998 - val_loss: 0.3384 - val_accuracy: 0.9000
Epoch 9/10
100/100 [=====] - 8s 81ms/step - loss: 0.3371 - ac
curacy: 0.9000 - val_loss: 0.3363 - val_accuracy: 0.9000
Epoch 10/10
100/100 [=====] - 5s 49ms/step - loss: 0.3355 - ac
curacy: 0.9000 - val_loss: 0.3352 - val_accuracy: 0.9000
```

### Saving the Model

```
In [27]: modelB.save('results/model-6-2-b')
```

```
INFO:tensorflow:Assets written to: results/model-6-2-b/assets
```

### Saving the Predictions

```
In [28]: predB = modelB.predict(validation_generator)

predB = pd.DataFrame(predB)
predB.to_csv('results/model-6-2-b/predictions.csv')
```

## Saving the Metrics

```
In [31]: test_lossB, test_accB = modelB.evaluate(validation_generator)

1563/1563 [=====] - 43s 27ms/step - loss: 0.3345 -
accuracy: 0.9000

In [32]: print("Test Accuracy: ", test_accB)
print("Test Loss: ", test_lossB)

Test Accuracy: 0.9000139236450195
Test Loss: 0.3345222771167755

In [33]: lines = ['Test Accuracy: 0.9000139236450195', 'Test Loss: 0.33452227711677

In [34]: with open('results/model-6-2-b/metrics.txt', 'w') as f:
f.write('\n'.join(lines))
```

## Saving the Validation Plots

```
In [39]: acc = historyB.history['accuracy']
val_acc = historyB.history['val_accuracy']

loss = historyB.history['loss']
val_loss = historyB.history['val_loss']

In [40]: epochs = range(1, len(acc) + 1)



In [41]: plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')

plt.legend()

plt.show()
```



```
In [42]: plt.savefig('results/model-6-2-b/validationplot.jpg')
```

<Figure size 432x288 with 0 Axes>